

Neobuffer: Cross-Process Channels in Rust

Bernardo Meurer

June 28th, 2019

Standard Cognition



- Bernardo Meurer
(`lovesegfault`)
- Systems Engineer @
Standard Cognition
<http://standard.ai>
- `bernardo@standard.ai`
- <http://lovesegfault.com/>

Outline

- Motivation
- Requirements
- Alternatives
- Neobuffer
 - Components
 - Ringbuffer
 - Sink
 - Source
 - Quirks
 - Results
 - Future work
- Questions


- Standard Cognition
- SF + Tokyo
- Autonomous Checkout
- Python + Rust
- Many interesting problems to solve
- We're hiring!



Motivation

- We deal with a many cameras.

Motivation



Motivation

- We deal with a many cameras.
- We do experiments very frequently.

Motivation



Motivation

- We deal with a many cameras.
- We do experiments very frequently.
- How to multiplex the video streams?

Motivation

- We deal with a many cameras.
- We do experiments very frequently.
- How to multiplex the video streams?
- How to unify the API for controlling distinct cameras?

- We still haven't quite decided on a name...

- We still haven't quite decided on a name...
- Camera daemon

- We still haven't quite decided on a name...
- Camera daemon
- Clients use ZMQ to communicate with it

- We still haven't quite decided on a name...
- Camera daemon
- Clients use ZMQ to communicate with it
- Has client libs for Rust, and Python (PyO3)


- We still haven't quite decided on a name...
- Camera daemon
- Clients use ZMQ to communicate with it
- Has client libs for Rust, and Python (PyO3)
- Reasonably large Rust project

Sensord/Camd

```
$ tokei ./sensord
```

| Language | Files | Lines | Code | Comments | Blanks |
|------------|-------|-------|------|----------|--------|
| C Header | 2 | 4 | 4 | 0 | 0 |
| GLSL | 2 | 63 | 54 | 1 | 8 |
| Markdown | 1 | 29 | 29 | 0 | 0 |
| Nix | 1 | 44 | 34 | 0 | 10 |
| Python | 1 | 28 | 17 | 4 | 7 |
| Rust | 47 | 9391 | 7457 | 880 | 1054 |
| Plain Text | 1 | 13 | 13 | 0 | 0 |
| TOML | 14 | 274 | 240 | 0 | 34 |
| Total | 69 | 9846 | 7848 | 885 | 1113 |

- We still haven't quite decided on a name...
- Camera daemon
- Clients use ZMQ to communicate with it
- Has client libs for Rust, and Python (PyO3)
- Reasonably large Rust project
- A small part of a huge system



- We still haven't quite decided on a name...
- Camera daemon
- Clients use ZMQ to communicate with it
- Has client libs for Rust, and Python (PyO3)
- Reasonably large Rust project
- A small part of a huge system
- One challenge was clear from inception...

How to connect **camd** to the client processes?

Channel X: Requirements

- Fast

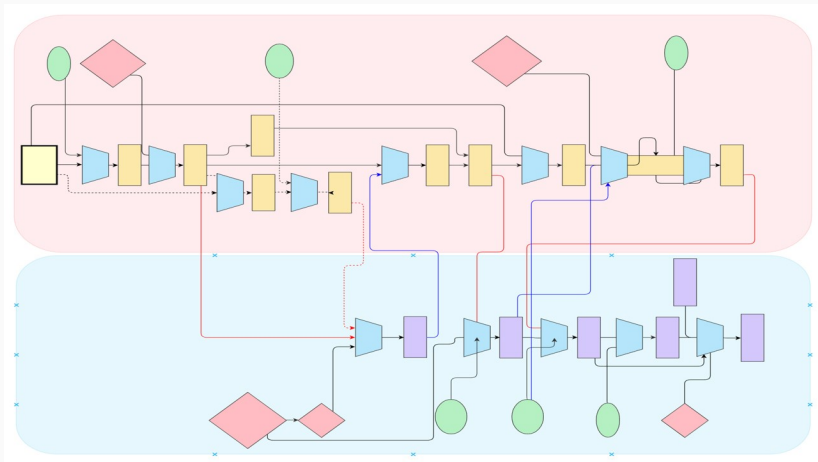
Channel X: Requirements

- Fast
- Lock-Free

Channel X: Requirements

- Fast
- Lock-Free
- Flexible

Channel X: Requirements



Channel X: Requirements

- Fast
- Lock-Free
- Flexible
- Cross-thread

Channel X: Requirements

- Fast
- Lock-Free
- Flexible
- Cross-thread
- Cross-process

Channel X: Requirements

- Fast
- Lock-Free
- Flexible
- Cross-thread
- Cross-process
- Safe

Channel X: Requirements

- Fast
- Lock-Free
- Flexible
- Cross-thread
- Cross-process
- Safe
- Multi-consumer

Channel X: Who could it be?

- stdlib channels?



Channel X: Who could it be?

- stdlib channels?
- crossbeam?



Channel X: Who could it be?

- stdlib channels?
- crossbeam?
- rb?



Channel X: Who could it be?

- `stdlib` channels?
- `crossbeam`?
- `rb`?
- `ringbuf`?





- Written from scratch
- 100% Rust
- Collaboration with **ekleog** (Leo Gaspard) and **nagisa** (Simonas Kaslauskas)
- Based on discussions with **eddyb** and **amanieu**

Neobuffer: Components

- Common traits for process-level **Sync** and **Send**
 - **interprocess-traits** — **ProcSync** and **ProcSend**
 - github.com/standard-ai/interprocess-traits

Neobuffer: Components

- Common traits for process-level **Sync** and **Send**
 - **interprocess-traits** — **ProcSync** and **ProcSend**
 - github.com/standard-ai/interprocess-traits
- Shared memory for cross-process data sharing
 - **shmem** — An ergonomic wrapper around shared memory for Rust
 - github.com/standard-ai/shmem

Neobuffer: Components

- Common traits for process-level **Sync** and **Send**
 - **interprocess-traits** — **ProcSync** and **ProcSend**
 - github.com/standard-ai/interprocess-traits
- Shared memory for cross-process data sharing
 - **shmem** — An ergonomic wrapper around shared memory for Rust
 - github.com/standard-ai/shmem
- Unix Domain Sockets for sending file descriptors across processes
 - **sendfd** — Utility crate that simplifies **fd** sharing.
 - github.com/standard-ai/sendfd

Neobuffer: Components

- Common traits for process-level **Sync** and **Send**
 - **interprocess-traits** — **ProcSync** and **ProcSend**
 - github.com/standard-ai/interprocess-traits
- Shared memory for cross-process data sharing
 - **shmem** — An ergonomic wrapper around shared memory for Rust
 - github.com/standard-ai/shmem
- Unix Domain Sockets for sending file descriptors across processes
 - **sendfd** — Utility crate that simplifies **fd** sharing.
 - github.com/standard-ai/sendfd
- Ringbuffer for the underlying data structure
 - **Atomics** for the indexes
 - Careful application of memory orderings to guarantee consistency
 - Safe wrap around
 - github.com/standard-ai/neobuffer

Neobuffer: Ringbuffer, Sink, and Source

```
/// A shared-memory ring buffer that contains items of type `T`, has one writer
/// and `NbReaders` readers that each read all the items the writer writes.
pub struct RingBuffer<T, NbReaders> {
    /// The shared memory region
    data: Shared<c_void>,

    /// The maximum number of elements this ring buffer can contain at a time.
    size: u64,

    /// Whether the one writer of this ring buffer has already been given out.
    writer_given: AtomicBool,

    /// The number of readers that have already been given out.
    readers_given: AtomicUsize,

    /// Phantom argument to keep type arguments alive.
    phantom: PhantomData<(T, NbReaders)>,
}
```

Neobuffer: Ringbuffer, Sink, and Source

- `Shared<T>` just means `T` is in shared memory


Neobuffer: Ringbuffer, Sink, and Source

- `Shared<T>` just means `T` is in shared memory
- `Shared<c_void>` is a special case for data not sized

Neobuffer: Ringbuffer, Sink, and Source

- `Shared<T>` just means `T` is in shared memory
- `Shared<c_void>` is a special case for data not sized
- But what exactly is in `Shared<c_void>`?

Neobuffer: Ringbuffer, Sink, and Source



Neobuffer: Ringbuffer, Sink, and Source

- `Shared<T>` just means `T` is in shared memory
- `Shared<c_void>` is a special case for data not sized
- But what exactly is in `Shared<c_void>`?
- `M` is the Ringbuffer metadata, and the magic smoke of Neobuffer

Neobuffer: Ringbuffer, Sink, and Source

```
/// The metadata for the data stored in the `Shared`.
///
/// The actual data in the ringbuffer precedes this metadata, so that it can be
/// page-aligned. Elements between the minimal element of `reader_steps` and
/// `writer_step` (modulo `Size`) are the only initialized ones. The safety of
/// this relies on the fact that things in a `Shared` are guaranteed to not be
/// dropped, and to the dropping operations being manually handled.
#[repr(C)]
struct Metadata {
    /// Number of readers allocated in the metadata area (following this struct).
    ///
    /// This is immutable.
    reader_count: usize,

    /// The number of items that have ever been written. Monotonically
    /// increasing. `u64` is used here because `u32` might not be enough to
    /// handle a reasonable number of ring buffer additions. However, all uses
    /// that are bounded by `Size` can safely assume it fits in `usize`.
    writer_step: AtomicU64,

    /// The number of outstanding references to this shared memory region.
    ///
    /// This knowledge is used to know when the remaining elements should be
    /// dropped.
    references: AtomicU64,
}
```

Neobuffer: Ringbuffer, Sink, and Source

```
impl<T, R> RingBuffer<T, R>
where
    // Ts are not created/inserted/shared here, so T does not need any bound here.
    R: Unsigned + NonZero,
{
    pub fn new(size: usize) -> Result<RingBuffer<T, R>, shmem::Error> {...}
    pub fn get_sink(&self) -> Result<Sink<T>, Error> {...}
    pub fn get_source(&self) -> Result<Source<T>, Error> {...}
}
```

Neobuffer: Ringbuffer, Sink, and Source

```
/// A sink that can send data into a `RingBuffer`.
pub struct Sink<T> {
    /// Shared pointer to the data.
    data: Shared<c_void>,

    /// Number of `T` that `data` has been allocated with.
    size: u64,

    /// Local copy of `data.writer_step`, that is not atomic for performance.
    writer_step: u64,

    /// Local helper that contains a value guaranteed to be lower to or equal
    /// to all values in `data.reader_steps`
    minimum_reader_step: Cell<u64>,

    /// Phantom argument to keep type arguments alive.
    phantom: PhantomData<T>,
}
```

Neobuffer: Ringbuffer, Sink, and Source

```
impl<T> Sink<T> {  
    pub fn push(&mut self, item: T) -> nb::Result<(), Error> {...}  
}
```


Neobuffer: Ringbuffer, Sink, and Source

```
/// A sink that can read data from a `RingBuffer`.
pub struct Source<T> {
    /// Shared pointer to the data.
    data: Shared<c_void>,

    /// Size of the ring buffer this is linked with.
    size: u64,

    /// Identifier of this reader in the `data.reader_steps` array.
    reader_id: usize,

    /// Local copy of `data.reader_steps[reader_id]`, for performance.
    reader_step: u64,

    /// Under-approximation of `data.writer_step`, local for performance.
    minimum_writer_step: Cell<u64>,

    /// Phantom argument to keep type arguments alive.
    phantom: PhantomData<T>,
}
```

Neobuffer: Ringbuffer, Sink, and Source

```
impl<T> Source<T> {  
    pub fn available_size(&self) -> usize {...}  
  
    /// Advance the reader discarding `n` lowest numbered elements in the buffer.  
    ///  
    /// This allows elements, which have been advanced past by all readers, to  
    /// be overwritten.  
    pub fn advance(&mut self, n: usize) -> nb::Result<(), Infallible> {...}  
  
    /// Get a reference to `i`-th element stored in the buffer.  
    ///  
    /// Once the use of the reference is done, [Source::advance`] must be  
    /// called to advance the reader.  
    pub fn get(&self, i: usize) -> nb::Result<&T, Infallible> {...}  
  
    /// Handy combination of get() and advance()  
    pub fn pop(&mut self) -> nb::Result<T, Infallible> {...}  
}
```

- Consumption is local

- Consumption is local
- Relative indexing

- Consumption is local
- Relative indexing
- Only works on Linux

- Consumption is local
- Relative indexing
- Only works on Linux
- You have to be surgically careful with atomic orderings

Neobuffer: Results

| | |
|--|--|
| one_writer_many_batching_readers_big | ... 10,693,940 ns/iter (+/- 931,225) = 1568 MB/s |
| one_writer_many_batching_readers_huge | ... 2,207,889 ns/iter (+/- 370,206) = 30395 MB/s |
| one_writer_many_batching_readers_small | ... 9,169,666 ns/iter (+/- 411,746) = 7 MB/s |
| one_writer_many_readers_big | ... 7,919,200 ns/iter (+/- 584,713) = 2118 MB/s |
| one_writer_many_readers_huge | ... 2,219,540 ns/iter (+/- 186,404) = 30235 MB/s |
| one_writer_many_readers_small | ... 6,706,252 ns/iter (+/- 480,979) = 9 MB/s |
| one_writer_one_batching_reader_big | ... 3,522,033 ns/iter (+/- 18,138) = 4763 MB/s |
| one_writer_one_batching_reader_huge | ... 2,188,906 ns/iter (+/- 19,540) = 30658 MB/s |
| one_writer_one_batching_reader_small | ... 3,369,997 ns/iter (+/- 314,502) = 19 MB/s |
| one_writer_one_reader_big | ... 5,138,619 ns/iter (+/- 263,518) = 3264 MB/s |
| one_writer_one_reader_huge | ... 2,196,294 ns/iter (+/- 31,967) = 30555 MB/s |
| one_writer_one_reader_small | ... 3,213,515 ns/iter (+/- 360,083) = 20 MB/s |
| one_writer_two_batching_readers_big | ... 5,827,360 ns/iter (+/- 160,344) = 2879 MB/s |
| one_writer_two_batching_readers_huge | ... 2,201,807 ns/iter (+/- 14,876) = 30478 MB/s |
| one_writer_two_batching_readers_small | ... 4,877,466 ns/iter (+/- 184,255) = 13 MB/s |
| one_writer_two_readers_big | ... 5,310,644 ns/iter (+/- 149,770) = 3159 MB/s |
| one_writer_two_readers_huge | ... 2,195,695 ns/iter (+/- 91,573) = 30563 MB/s |
| one_writer_two_readers_small | ... 4,028,670 ns/iter (+/- 403,310) = 16 MB/s |

Assuming 4 parallel readers and a buffer size of 8:

- 7 Mpps (1 byte packets)
 - 50Mpps with a buffer size of 2048

Assuming 4 parallel readers and a buffer size of 8:

- 7 Mpps (1 byte packets)
 - 50Mpps with a buffer size of 2048
- 300 Gbps (2MB packets)

- Improve ergonomics around failure cases

Neobuffer: Future work

- Improve ergonomics around failure cases
- Split pushing into allocate + commit (remove a copy/clone)

Neobuffer: Future work

- Improve ergonomics around failure cases
- Split pushing into allocate + commit (remove a copy/clone)
- Swappable backend

Neobuffer: Future work

- Improve ergonomics around failure cases
- Split pushing into allocate + commit (remove a copy/clone)
- Swappable backend
- MPMC

Neobuffer: Future work

- Improve ergonomics around failure cases
- Split pushing into allocate + commit (remove a copy/clone)
- Swappable backend
- MPMC
- Futures

Neobuffer: Future work

- Improve ergonomics around failure cases
- Split pushing into allocate + commit (remove a copy/clone)
- Swappable backend
- MPMC
- Futures
- Global consumption

Summary

- Lock-free, fast, cross-process channel

Summary

- Lock-free, fast, cross-process channel
- Ringbuffer in shared memory

Summary

- Lock-free, fast, cross-process channel
- Ringbuffer in shared memory
- Public supporting libraries

Summary

- Lock-free, fast, cross-process channel
- Ringbuffer in shared memory
- Public supporting libraries
- Atomics to manage read/write indexes

Summary

- Lock-free, fast, cross-process channel
- Ringbuffer in shared memory
- Public supporting libraries
- Atomics to manage read/write indexes
- Safe interface for cross-process communication

There will be no questions

Are there any questions?

This presentation was made using Free (as in Freedom) and Open Source software:

- Vim
- \LaTeX + Beamer + **metropolis**
- Graphviz
- Inkscape
- Some dude on Fiverr